

“Structured Query Language”

- It is NOT much like other programming languages
- It is NOT PROCEDURAL
- It does not process one record at a time, rather, it is a SET processing language
- All inputs to SQL are tables
- The output from a query is a table
- Output from a query referred to as the “Answer Set”
- Some queries may produce “interim” temporary answer sets

“Structured Query Language”

- It is a relatively simple language – brief syntax, few commands
- It is a relatively powerful language – a FEW lines of code can accomplish a LOT of work
- ANSI (American National Standards Institute) maintains a specification for “standard” SQL
- Each DBMS manufacturer follows the ANSI standard, but also adds extended features unique to their SQL

Useful Tool for Managing your Databases

- “SQLYog” from WebYog
- The community edition is free
- <https://github.com/webbyog/sqllyog-community/wiki/Downloads>

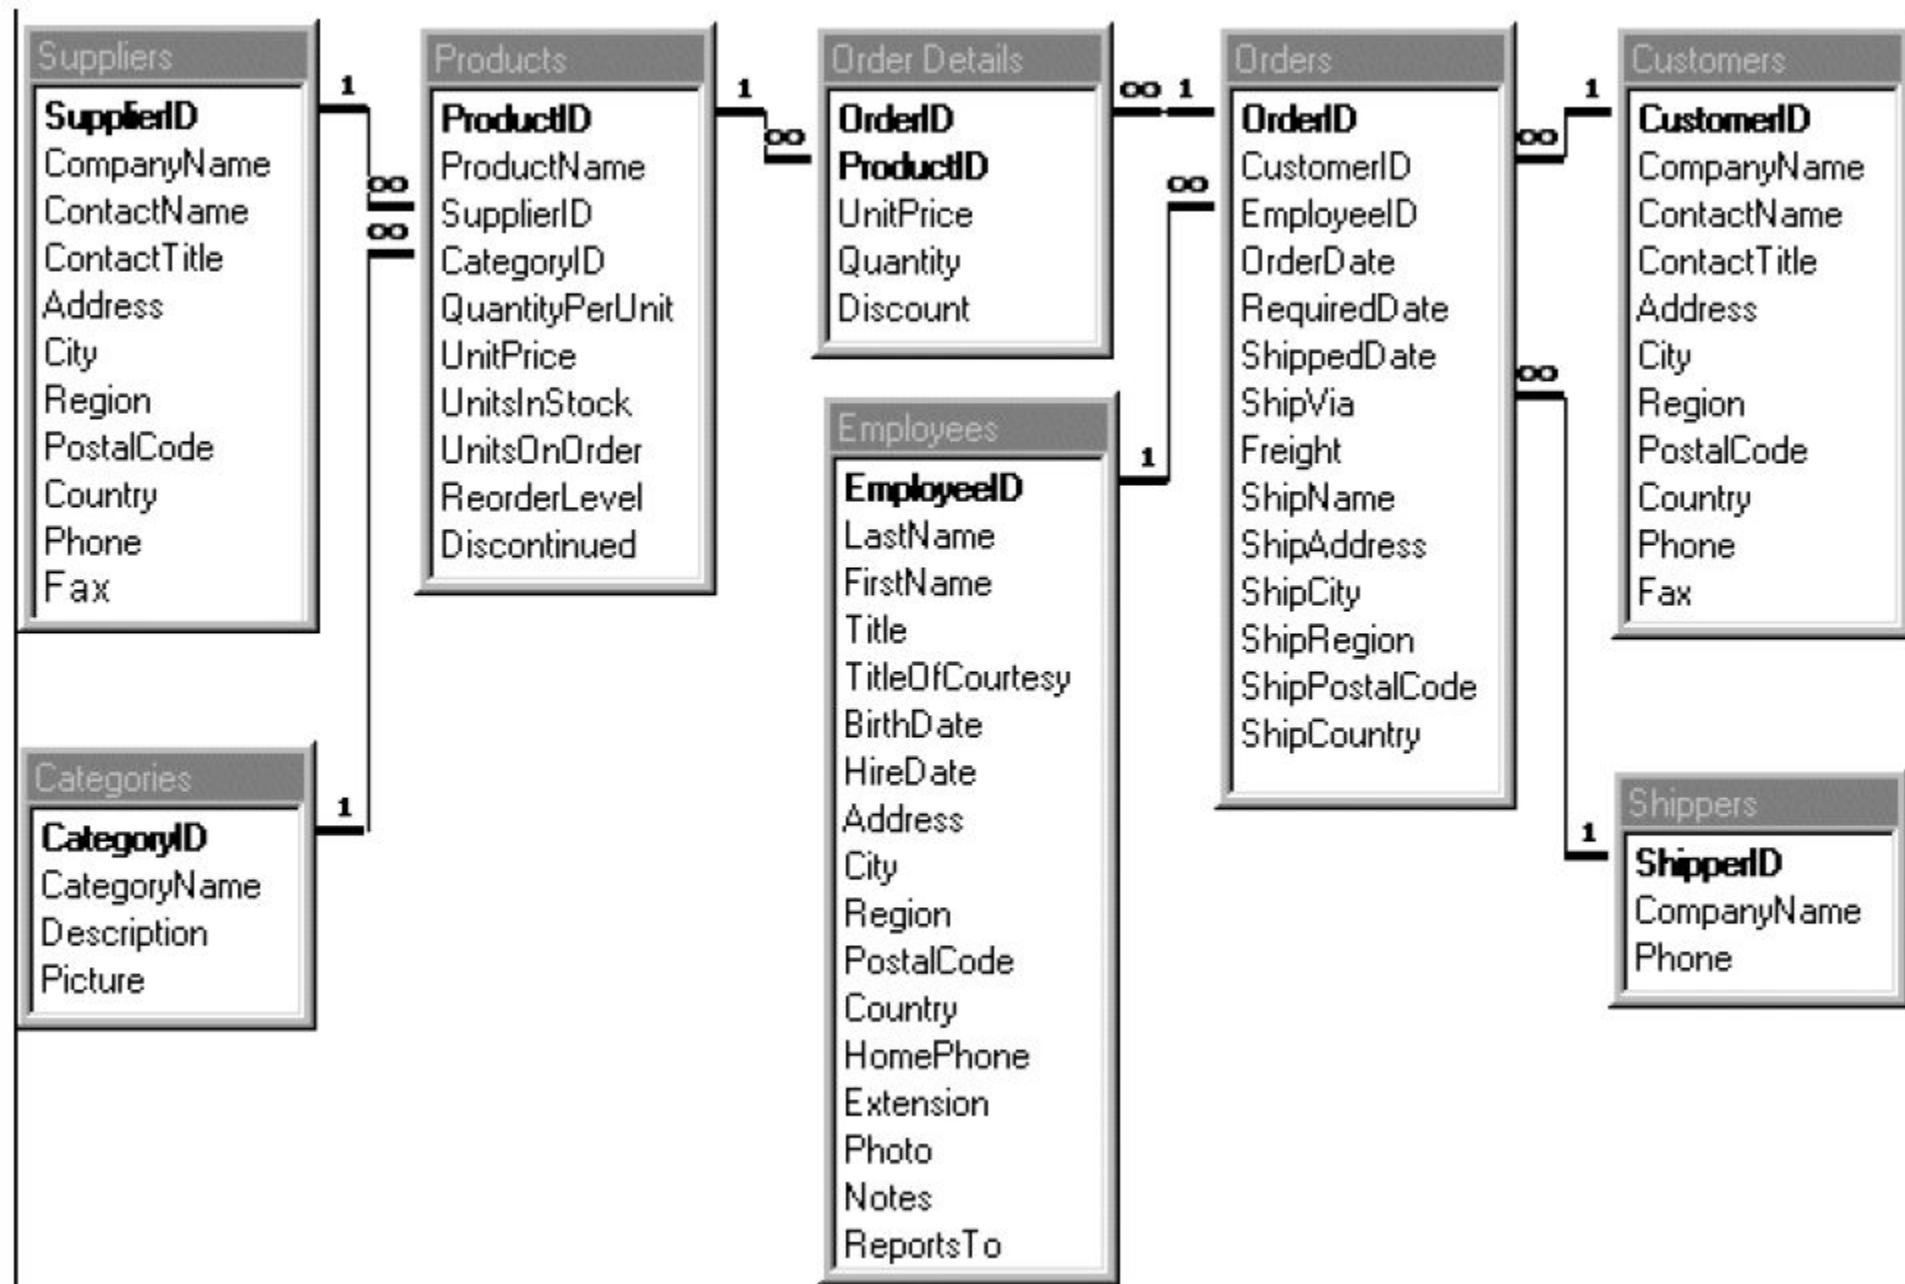
USE statement

```
USE <database>;
```

- Tells the Query Engine which database you want to use for your query

SELECT statement

```
SELECT <column1>, <column2>, <column3>,  
<literal>, <math expression>  
    FROM <table A> ;
```



Examples

```
select *  
    from nwEmployees;
```

```
use NorthWinds;
```

```
select EmployeeID, LastName, FirstName  
    from nwEmployees;
```

SELECT statement

- Literals may be either 'Character' (in quotes) or Numeric

- Math expressions

Only use with columns defined as numeric data types

+ Add

- Subtract

* Multiply

/ Divide

** Exponent

SELECT statement

- Rename a column in the answer set with “AS”
`Select employeeID as 'EMP'`
- Concatenate character columns with (MySQL only)
`Concat (<column1>, column2>)`
- Comment out a line or part of a line of code by prefixing it with “- -” or embedding a “#”
- Limit the size of the answer set with “limit” (MySQL)
`Select LastName, FirstName
from nwEmployees limit 100;`

Examples

```
select 'Roster', LastName, FirstName  
    from nwEmployees;
```

```
select 'Roster' as 'Type', LastName, FirstName  
    from nwEmployees;
```

```
select 22, LastName, FirstName  
    from nwEmployees;
```

```
select 2 * 2, LastName, FirstName  
    from nwEmployees;
```

```
select concat(FirstName, ' ', LastName)  
    from nwEmployees;
```

SELECT statement

```
SELECT <column1>, <column2>, <column3>,  
<literal>, <math expression> AS <label>  
    FROM    <table A>  
    WHERE   <condition> ;
```

- The WHERE clause results in a subset of ROWs to appear in the answer set
- The condition in the WHERE clause takes this format:

< operand > < operator > < operand >

- Operands may be columns or literals or expressions
- Operator may be

=	Equals	Like
<>	Not equals	Between
>	Greater than	In
<	Less than	

- Operator may be: In or Like
In (literal, literal, literal)
Like 'string' with % or _ as a wildcard
- Multiple conditions may be joined with Boolean operators
AND, OR
- Conditions may be negated with Boolean operator
NOT
- Answer Set rows may be sorted with “Order By”
- Order By defaults to Ascending, can specify DESC

Distinct:

- The answer set may contain duplicate rows
- The “`distinct`” keyword before a column removes duplicates
- Example:
 - 87 Customers, each one has a country
 - How many distinct countries are they from?

Examples

```
select CustomerID, ContactName, Region, Country
    from nwCustomers;
```

```
select CustomerID, ContactName, Region, Country
    from nwCustomers
    where Country = 'Brazil';
```

```
select CustomerID, ContactName, Region, Country
    from nwCustomers
    where Country <> 'Brazil';
```

```
select ProductID, ProductName, UnitPrice
    from nwProducts
    where UnitPrice > 60;
```

Examples

```
select ProductID, ProductName, UnitPrice
  from nwProducts
 where UnitPrice between 20 and 30;
```

```
select ProductID, ProductName, categoryid, UnitPrice
  from nwProducts
 where UnitPrice between 20 and 30
    and categoryID in (2, 4, 6);
```

```
select ProductID, ProductName, QuantityPerUnit
  from nwProducts
 where QuantityPerUnit like '%jars%';
```


Examples: Using distinct

```
Select CompanyName, ContactName, Country  
    from nwCustomers;
```

```
Select Country  
    from nwCustomers;
```

```
Select Distinct Country  
    from nwCustomers;
```

Handling Dates in MySQL

- MySQL supports DATE, DATETIME, and TIMESTAMP data types
- Columns with a data type of “TIMESTAMP” are stored as a 4-byte binary integer representing the number of seconds since 1970-01-01 00:00:00 UTC. TIMESTAMP has a range of '1970-01-01 00:00:01' UTC to '2038-01-19 03:14:07' UTC.
- If no value is provided for the TIME portion of a DATETIME column, it defaults to 00:00.00.0000

Handling Dates in MySQL

- To make it easier for humans to deal with date/time, MySQL allows us to reference dates/times in this format:
 - YYYY-MM-DD and HH:MM.SS.nnn
- If you pass the date to MySQL as text in YYYY-MM-DD format, it will automatically convert it to the proper binary number
- If you pass the time to MySQL as text in HH:MM.SS.nnn format, it will automatically convert it to the proper binary number

YYYY-MM-DD and hh:mm:ss.nnn

- **YYYY** is **four digits** from 1000 through 9999 that represent a year.
- **MM** is **two digits**, ranging from 01 to 12, that represent a month in the specified year.
- **DD** is **two digits**, ranging from 01 to 31 depending on the month, that represent a day of the specified month.
- **hh** is **two digits**, ranging from 00 to 23, that represent the hour.
- **mm** is **two digits**, ranging from 00 to 59, that represent the minute.
- **ss** is **two digits**, ranging from 00 to 59, that represent the second.
- **nnn** is **zero to three digits**, ranging from 0 to 999, that represent the fractional seconds.

Examples: Using DATES

```
Select Now();
```

```
Select Curdate();
```

```
Select Curtime();
```

```
Select Lastname, Firstname, Extract(Year From HireDate) AS HireYear  
FROM NWEmployees;
```

```
SELECT EmployeeID, Lastname, Firstname,  
ROUND(DATEDIFF(HireDate, BirthDate)/365,0) AS HIRE_AGE  
FROM NWEmployees;
```

Examples: Using DATES

```
SELECT DATE_FORMAT(HireDate, '%b %d %Y %h:%i %p') FROM nwEmployees;
```

Format	Description
%a	Abbreviated weekday name (Sun-Sat)
%b	Abbreviated month name (Jan-Dec)
%c	Month, numeric (0-12)
%D	Day of month with English suffix (0th, 1st, 2nd, 3rd, 4th, 5th, 6th, 7th, 8th, 9th, 10th, 11th, 12th, 13th, 14th, 15th, 16th, 17th, 18th, 19th, 20th, 21st, 22nd, 23rd, 24th, 25th, 26th, 27th, 28th, 29th, 30th, 31st)
%d	Day of month, numeric (00-31)
%e	Day of month, numeric (0-31)
%f	Microseconds (000000-999999)
%H	Hour (00-23)
%h	Hour (01-12)
%l	Hour (01-12)
%i	Minutes, numeric (00-59)
%j	Day of year (001-366)
%k	Hour (0-23)
%l	Hour (1-12)

SELECT statement

```
SELECT <column1>, <column2>, <column3>, <literal>,  
<math expression>  
    FROM <tableA>  
    WHERE <condition>  
ORDER BY <column1>, <column2> [DESC] ;
```

SQL provides the following GROUP FUNCTIONS

SUM - Provides the sum of the values in a column across many rows

AVG - Provides the average of the values in a column across many rows

COUNT - Provides a count of how many rows have a value in a column, counted across many rows

MIN - Provides the lowest value in a column across many rows

MAX - Provides the highest value in a column across many rows

GROUP FUNCTIONS

- SUM, AVG must only be used with NUMERIC columns
- MIN, MAX can be used with any data type
- COUNT can be used with any column, or with a (*) to simply count rows
- **Group functions require SQL to create an interim answer set, and then process the group function against the interim answer set, delivering a final answer set that contains only the final total for the function. Always returns an integer value.**
- When you combine a GROUP FUNCTION with a WHERE clause, keep in mind that the **WHERE clause simply reduces the number of rows in the INTERIM answer set** before the GROUP function does its calculation.

Examples

```
select COUNT(*) as 'Total'
      from nwEmployees;
```

```
select COUNT(Distinct Country) as 'Countries'
      from nwCustomers;
```

```
select SUM(UnitPrice) as 'Total Price'
      from nwProducts;
```

```
select MAX(UnitPrice) as 'High Price'
      from nwProducts;
```

```
select MIN(UnitPrice) as 'Low Price'
      from nwProducts
      where UnitPrice > 0;
```

```
select AVG(UnitsInStock) as 'Average Inventory'
      from nwProducts;
```

GROUP BY

- Group functions process against an interim answer set to return a value across many rows.
- Using a GROUP BY clause enables SQL to provide subtotals. The GROUP BY tells SQL to **perform the group function against a subset of rows in the interim answer set** and provide a total for each subset of rows.
- **When using a GROUP BY every column in the SELECT statement must either be a GROUP FUNCTION or a COLUMN that you are grouping by.**

Examples

```
select Country, COUNT(*) as 'Total'  
  from nwCustomers  
 GROUP BY Country;
```

```
select Country, COUNT(Country) as 'Total'  
  from nwCustomers  
 GROUP BY Country;
```

Why are these the same?

```
select CustomerID, Country, COUNT(Country) as 'Total'  
  from nwCustomers  
 GROUP BY Country;
```

Why is this incorrect?

More Examples

From which supplier does Northwinds carry the most inventory?

```
select SupplierID, SUM(UnitsInStock) as 'Inventory'  
from nwProducts  
group by SupplierID;
```

In which month in 2013 did Northwinds ship the most orders?

```
SELECT EXTRACT(MONTH FROM ShippedDate) AS 'Month',  
COUNT(OrderID) AS 'Orders'  
FROM nwOrders  
WHERE EXTRACT(YEAR FROM ShippedDate) = '2013'  
GROUP BY EXTRACT(MONTH FROM ShippedDate)  
ORDER BY 2 DESC  
Limit 1;
```

HAVING

- Is simply like a WHERE clause against the answer set when you use a GROUP BY

Examples

```
select Country, COUNT(*) as 'Total'
  from nwCustomers
 GROUP BY Country;
```

In which countries does Northwinds have more than five customers?

```
select Country, COUNT(*) as 'Total'
  from nwCustomers
 GROUP BY Country
 HAVING COUNT(*) > 5
 order by 2 desc;
```

SubQuery

Simply: a query within a query. The answer set to an “inner” query is used as a predicate in a where clause in the “outer” query.

- The subquery must **return only one column**.
- If the outer query WHERE clause contains an “**equals**” condition, the subquery must return **ONE row**.
- If the outer query WHERE clause contains an “**in**” condition, the subquery may return **multiple rows**, presented as a list of values.
- The Subquery is embedded within parentheses
- Outer and inner queries can hit two different tables

Examples

```
select ProductID, ProductName, UnitPrice
  from nwProducts
 where UnitPrice = (
    select MAX(UnitPrice)
      from nwProducts );
```

Note that with the “equals” condition, the inner query returns only one value (one row, one column)

Examples

```
select CustomerID, OrderID
  from nwOrders
 where OrderID in (
    select OrderID
      from nwOrderDetails
     where Quantity > 100 )
 Order by CustomerID;
```

Note that with the “in” condition, the inner query returns many values (many rows, one column) as a list

Uses Two different tables

Getting Data from Multiple Tables – The JOIN

- In order to run a query that retrieves data from multiple tables, those tables must be **JOINED**.
- Joining two tables requires that the two tables have a common key (typically a foreign key relationship) that appears in both tables.
- The common key columns need NOT have the same name, but must be of the **same data type and length**.
- A JOIN is one of the most **resource intensive** activities one can do in a relational database.

Basic Example

- Let's join nwOrders to nwEmployees
- nwOrders has 830 rows, each with an EmployeeID
- nwEmployees has 9 rows, each with an EmployeeID
- They have a common key: EmployeeID (primary key in nwEmployees; foreign key in nwOrders)
- We want SQL to join the rows in nwEmployees and nwOrders where the EmployeeID matches

Provide a listing showing Northwinds employees, sorted by LastName, and a count of each employee's orders

```
Select LastName, Firstname, count (OrderID) as 'Orders'
from nwEmployees, nwOrders
where nwEmployees.EmployeeID =
      nwOrders.EmployeeID
GROUP BY LastName, FirstName
Order By 1;
```

Qualifying the Column Names

- Since the column “EmployeeID” exists in BOTH tables in this query, when referring to EmployeeID, we need to tell SQL which one.
- Therefore, we suffix the table name in front of the column name separated by a “ . ”
- Failure to fully qualify the column name will result in an “ambiguous column” error

Alternative

To save some typing, we can define an “alias” for each table. We can temporarily – only for the duration of this query -- rename the nwEmployees table “E”, and rename the nwOrders table “O”.

```
Select LastName, Firstname, count(OrderID) as 'Orders'
  from nwEmployees E, nwOrders O
 where E.EmployeeID = O.EmployeeID
GROUP BY LastName, FirstName
Order By 1;
```

Alternative

- SQL allows another syntax option for doing the JOIN.
- These queries are equivalent:

```
Select LastName, Firstname, count(OrderID) as 'Orders'
  from nwEmployees E, nwOrders O
  where E.EmployeeID = O.EmployeeID
 GROUP BY LastName, FirstName
 Order By 1;
```

```
Select LastName, Firstname, count(OrderID) as 'Orders'
  from nwEmployees E JOIN nwOrders O
  on E.EmployeeID = O.EmployeeID
 GROUP BY LastName, FirstName
 Order By 1;
```


Beware the Cartesian Product

- Product = one table multiplied by another table
- **The JOIN often creates a product, then selects rows from the product where the keys match**
- For example, let's join nwOrders to nwEmployees
- nwOrders has 14 columns, 830 rows
- nwEmployees has 17 columns, 9 rows
- The Cartesian product has 31 (14+17) columns, and 7470 (830 * 9) rows – most of which are meaningless

Cartesian Product

- SQL must go through the Cartesian Product (which is an INTERIM answer set) row-by-row, and select only those rows where the EmployeeID from nwEmployees is equal to the EmployeeID from nwOrders
- We must include the WHERE clause that describes the selection
- Without WHERE a JOIN operation matches all necessary keys and will cause your answer set to include part or all of the Cartesian Product
- The JOIN requires SQL to do a lot of work which consumes a lot of disk I/O and memory (= expensive)

Cartesian Product from an unqualified join:

customerid	companyname	contactname	country	OrderID	CustomerID	Orderdate	shipcountry
GREAL	Great Lakes Food Market	Howard Snyder	USA	10262	RATTC	2013-07-22	USA
HUNGC	Hungry Coyote Import Store	Yoshi Latimer	USA	10262	RATTC	2013-07-22	USA
LAZYK	Lazy K Kountry Store	John Steel	USA	10262	RATTC	2013-07-22	USA
LETSS	Lets Stop N Shop	Jaime Yorres	USA	10262	RATTC	2013-07-22	USA
LONEP	Lonesome Pine Restaurant	Fran Wilson	USA	10262	RATTC	2013-07-22	USA
OLDWO	Old World Delicatessen	Rene Phillips	USA	10262	RATTC	2013-07-22	USA
RATTC	Rattlesnake Canyon Grocery	Paula Wilson	USA	10262	RATTC	2013-07-22	USA
SAVEA	Save-a-lot Markets	Jose Pavarotti	USA	10262	RATTC	2013-07-22	USA
SPLIR	Split Rail Beer & Ale	Art Braunschweiger	USA	10262	RATTC	2013-07-22	USA
THEBI	The Big Cheese	Liz Nixon	USA	10262	RATTC	2013-07-22	USA
THECR	The Cracker Box	Liu Wong	USA	10262	RATTC	2013-07-22	USA
TRAIH	Trails Head Gourmet Provisioners	Helvetius Nagy	USA	10262	RATTC	2013-07-22	USA
WHITC	White Clover Markets	Karl Jablonski	USA	10262	RATTC	2013-07-22	USA
GREAL	Great Lakes Food Market	Howard Snyder	USA	10269	WHITC	2013-07-31	USA
HUNGC	Hungry Coyote Import Store	Yoshi Latimer	USA	10269	WHITC	2013-07-31	USA
LAZYK	Lazy K Kountry Store	John Steel	USA	10269	WHITC	2013-07-31	USA
LETSS	Lets Stop N Shop	Jaime Yorres	USA	10269	WHITC	2013-07-31	USA
LONEP	Lonesome Pine Restaurant	Fran Wilson	USA	10269	WHITC	2013-07-31	USA
OLDWO	Old World Delicatessen	Rene Phillips	USA	10269	WHITC	2013-07-31	USA
RATTC	Rattlesnake Canyon Grocery	Paula Wilson	USA	10269	WHITC	2013-07-31	USA
SAVEA	Save-a-lot Markets	Jose Pavarotti	USA	10269	WHITC	2013-07-31	USA
SPLIR	Split Rail Beer & Ale	Art Braunschweiger	USA	10269	WHITC	2013-07-31	USA

Joining three or more tables

- Every PAIR of tables being joined must have a common key
- Every PAIR of common keys must have a condition stated in a WHERE clause or in the “ON” clause of the JOIN
- Otherwise, your JOIN is not fully qualified and will result in a Cartesian Product (meaningless output)

Examples – Joining three tables

Create a report showing each employee and the total value of their orders sorted from highest value to lowest. (Order Value = UnitPrice * Quantity for each item on the order.)

```
Select LastName, Firstname,  
       sum(UnitPrice * Quantity) as 'OrderValue'  
from   nwEmployees E  
JOIN   nwOrders O on E.EmployeeID = O.EmployeeID  
JOIN   nwOrderDetails D on O.OrderID = D.OrderID  
GROUP BY LastName, FirstName  
Order By 3 desc
```

Examples – Joining three tables

Same Query, Different Syntax

```
Select LastName, Firstname,  
       sum(UnitPrice * Quantity) as 'OrderValue'  
from nwEmployees E, nwOrders O, nwOrderDetails D  
where E.EmployeeID = O.EmployeeID  
      and O.OrderID = D.OrderID  
GROUP BY LastName, FirstName  
Order By 3 desc
```

Inner join ONLY includes matching rows

Explicit inner join

```
SELECT * FROM employee  
INNER JOIN department ON employee.DepartmentID =  
    department.DepartmentID
```

Implicit inner join:

```
SELECT * FROM employee, department  
WHERE employee.DepartmentID =  
    department.DepartmentID
```

LEFT outer join

```
SELECT * FROM employee  
LEFT OUTER JOIN department ON  
employee.DepartmentID =  
department.DepartmentID
```

Returns ALL rows from LEFT table and only matching rows from RIGHT table.

RIGHT outer join

```
SELECT * FROM employee  
    RIGHT OUTER JOIN department ON  
employee.DepartmentID =  
department.DepartmentID
```

Returns ALL rows from RIGHT table and only matching rows from LEFT table.

Analysis using Outer Joins

- Are there some rows in `nworders` that do not have a valid foreign key reference to `nwcustomers`?
- Are there some customers that have no orders?

```
SELECT COUNT(customerid) FROM nwcustomers
```

- there are 87 customers in `nwcustomers`

```
SELECT COUNT(distinct customerid) FROM nworders
```

- there are 89 distinct customers in `nworders`

- What's the difference?

We want to find the orders in nworders whose customerID is NOT in nwcustomers

Method One: use a subquery

```
SELECT DISTINCT customerID
FROM nworders
WHERE customerID NOT IN (
    SELECT customerID FROM nwcustomers));
```

Method Two: use an outer join

```
SELECT DISTINCT O.customerID
FROM nworders O LEFT OUTER JOIN nwcustomers C
ON O.customerID = C.customerID
WHERE C.customerID IS NULL
```

This shows us FOUR customers who have orders in nworders that have no matching row in nwcustomers

We want to find the customers in nwcustomers who have no orders in nworders

Method One: use a subquery

```
SELECT customerID
FROM nwcustomers
WHERE customerID NOT IN (
    SELECT DISTINCT customerID FROM nworders);
```

Method Two: use an outer join

```
SELECT DISTINCT C.customerID
FROM nworders O RIGHT OUTER JOIN nwcustomers C
ON O.customerID = C.customerID
WHERE O.customerID IS NULL
```

This shows us TWO customers who have no orders in nworders

DDL = Data Definition Language

Some SQL commands are used to DEFINE or MODIFY the structures in the database.

Create

Alter

Drop

DML = Data Manipulation Language

Some SQL commands are used to MODIFY the data in the database

Update

Insert

Delete

CREATE statement

```
CREATE TABLE <table name>
(column    DATATYPE (L) ,
column    DATATYPE (L)  NOT NULL,
column    DATATYPE (L)  CONSTRAINT <constr
name> TYPE,
column    DATATYPE (L) )
TABLESPACE <tablespace name>;
```

ALTER statement

```
ALTER TABLE <table name>
```

```
    ADD COLUMN column DATATYPE (L) ,  
    CONSTRAINT <constr name> TYPE
```

```
ALTER TABLE <table name>
```

```
    DROP CONSTRAINT <constr name>
```

DROP statement

```
DROP TABLE <table name>
```

UPDATE statement

```
UPDATE <table name>  
SET column = <value>  
WHERE <condition>
```

NOTE: if the WHERE is omitted, ALL rows are updated.

DELETE statement

```
DELETE FROM <table name>  
WHERE <condition>
```

NOTE: if the WHERE is omitted, ALL rows are deleted.

INSERT statement

```
INSERT INTO <table name>  
VALUES (value, value, value, value);
```

(must have a value or NULL for every column in the table)

```
INSERT INTO <table name> (column, column,  
    column)  
VALUES (value, value, value);
```

(if no column/value is specified, NULL or default will be assigned)

Getting more help with SQL

Best Online Tutorial: <http://www.w3schools.com/sql/>